

Light-Weight Contexts: An OS Abstraction for Safety and Performance

James Litton, University of Maryland, College Park and Max Planck Institute for Software Systems (MPI-SWS); Anjo Vahldiek-Oberwagner, Eslam Elnikety, and Deepak Garg, Max Planck Institute for Software Systems (MPI-SWS); Bhattacharjee, University of Maryland, College Park; Peter Druschel, Max Planck Institute for Software Systems (MPI-SWS)

Introduction

- ▶ New OS Abstraction - light-weight contexts (*lwc*s)
- ▶ Light-Weight Contexts (LWC) is an operating system (OS) abstraction used for managing isolated contexts within a single process.
- ▶ Why *lwc*s
 - ▶ existing methods for session isolation are often slower than *lwc*s.
 - ▶ *lwc*-supported sensitive data compartments have negligible overhead on production servers.

Advantages

- ▶ /wCs enable a range of new in-process capabilities, such as
 - ▶ fast roll-back
 - ▶ protection rings (by credential restriction)
 - ▶ session isolation
 - ▶ Efficiency
 - ▶ Ease of use
 - ▶ protected compartments (using VM and resource mappings).

How does it help with security?

- ▶ It helps us achieve security with the following properties
 - ▶ Isolation
 - ▶ Sandboxing
 - ▶ Resource Control
 - ▶ Reduced Attack Surface
 - ▶ Improved Code Quality

Related Work

- ▶ Some other solutions include
 - ▶ Shreds [9] – however Lcw are fully independent of threads, require no compiler support, and rely on page-based hardware protection only. *lwCs* also provide protection rings and snapshots, which shreds do not.
 - ▶ Dune [4] – however it has higher overhead cost due to TLB misses and kernel calls.
 - ▶ Software fault isolation (SFI) [29] and NaCl [35] – *lwCs* instead allow fine-grained control over memory, file descriptors and other process credentials, and provide snapshots as part of an OS abstraction.
 - ▶ A few more including SpaceJMP [12], Corey [6], however they don't provide OS Snapshots and in-process isolation.

Why LcWs?

- ▶ While related work is being continued LcWs provide the following advantages
 - ▶ Lcws are fully independent of threads
 - ▶ Provide protection rings (Security) and snapshots
 - ▶ Low or negligible overhead cost
 - ▶ Fine-grained control over memory,
 - ▶ In-process isolation

Creating LcWs

- ▶ Starts with creating an LcW.
- ▶ The `lwCreate` call creates a new (child) `lwC` in the current process.
- ▶ the child `lwC`'s initial state is an identical copy of the calling (parent) `lwC`'s state, except for its descriptor.
- ▶ By default, the new `lwC` gets a private copy of the calling `lwC`'s state at the time of the call
- ▶ Shared memory regions in the calling `lwC` are shared with the new `lwC`.
- ▶ The implementation does not stop other threads executing in the parent `lwC` during an `lwCreate`.

Switching between *lwCs*

- ▶ The `lwSwitch` operation switches the calling thread to the *lwC* with descriptor *target*, passing *args* as parameters.
- ▶ `lwSwitch` retains the state of the calling thread in the present *lwC*.

How does it achieve Isolation?

- ▶ */wCs* do not have access to the state of each others' memory, file descriptors, and capabilities unless explicitly shared, they can provide strong isolation and privilege separation within a process.
- ▶ */wCs* can reliably prevent accidental leakage of private information across user sessions, isolate authentication credentials and other secrets
- ▶ An application that wishes to limit information flow across */wCs* should create */wCs* without the `LWC_SHARESIGNALS` option (the default).

Snapshot and rollback

Algorithm 1 Snapshot and rollback

```
1: function SNAPSHOT()
2:   new, caller, arg = lwCreate(default_spec, ...)
3:   if caller = -1 then                                ▷ parent
4:     return new
5:   else
6:     close(caller)
7:     return snapshot()
8: function ROLLBACK(snap)                                ▷ never returns
9:   lwSwitch(snap, 0)
10: function MAIN()
11:   ...                                                    ▷ initialize state
12:   snap = snapshot()
13:   ...                                                    ▷ serve request
14:   rollback(snap)
    ▷ kills current lwC, continues at line 12 in snap
```

lwC Implementation

- ▶ Like a process, each *lwC* has a file table, virtual memory space, and credentials associated with it.
- ▶ Memory
 - ▶ *lwCreate* replicates the vm space associated with the parent *lwC* in exactly the same manner as *fork*.
- ▶ File Table
 - ▶ By default, during a call to *lwCreate* all file descriptors are copied into the *lwC* file table in the same manner as *fork* except that any associated file descriptor overlay rights are copied as well
- ▶ Permissions and Overlays
 - ▶ An executing *lwC* interacts with another *lwC* within a process by either switching to it or by overlaying (some of) that *lwC*'s resources.

Evaluation

- ▶ The table compares the time to execute a lwSwitch

<i>lwC</i>	process	k-thread	u-thread
2.01 (0.03)	4.25 (0.86)	4.12 (0.98)	1.71 (0.06)

Conclusion

- ▶ *lwCs* provide isolation and privilege separation among program components within a process
- ▶ Also provides fast OS-level snapshots and coroutine style control transfer among contexts
- ▶ Provides fast roll-back

Reference

- ▶ ABADI, M., BUDIU, M., ERLINGSSON, U., AND LIGATTI, J. Control-flow integrity. In *Proceedings of the 12th ACM Conference on Computer and Communications Security (CCS)* (2005), pp. 340–353.
- ▶ [2] AVIRAM, A., WENG, S.-C., HU, S., AND FORD, B. Efficient system-enforced deterministic parallelism. In *Proceedings of the 9th USENIX Conference on Operating Systems Design and Implementation* (Berkeley, CA, USA, 2010), OSDI'10, USENIX Association, pp. 193–206.
- ▶ BANGA, G., DRUSCHEL, P., AND MOGUL, J. C. Resource containers: A new facility for resource management in server systems. In *Proceedings of the Third Symposium on Operating Systems Design and Implementation* (Berkeley, CA, USA, 1999), OSDI '99, USENIX Association, pp. 45–58.