



CS5375 Computer Systems Organization and Architecture

Lecture 18

Guest Instructors:

Ghazanfar Ali, Ghazanfar.Ali@ttu.edu

Mert Side, Mert.Side@ttu.edu

Department of Computer Science

Texas Tech University

Outline

- Programming Project #2 Walkthrough
- Midterm Exam Review

Reminder: Matrix Multiplication

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \times \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}$$

A B C

A, B and C are square matrices of size $N \times N$

a, b, c and d are submatrices of A, of size $N/2 \times N/2$

e, f, g and h are submatrices of B, of size $N/2 \times N/2$

```

// ----- CPUmatmul
void CPUmatmul(int N, double *x, double *y, double *ans)
{
    for(int i=0; i < N; i++) {
        for(int j=0; j < N; j++) {
            for(int k=0; k < N; k++) {
                ans[i*N+j] += (x[i*N+k]*y[k*N+j]);
            }
        }
    }
}

// ----- check
bool check(int N, double *ans)
{
    for(int i=0; i < N; i++) {
        for(int j=0; j < N; j++) {
            if(ans[i*N+j]!=20.0)return false;
        }
    }
    return true;
}

// ----- MAIN
int main(void)
{
    //size of matrix
    int N = 1<<9;
    int iter = 3;
    clock_t t;

    // Allocate Memory - accessible from CPU
    double *x = new double[N*N];
    double *y = new double[N*N];
    double *ans = new double[N*N];

    // .....
    // initialize x,y and ans arrays on the host
    for (int i = 0; i < N; i++) {
        for(int j=0; j < N; j++) {
            x[i*N+j]=5;
            y[i*N+j]=(i==j?1:0);
            ans[i*N+j]=(double)0.000000000000;
        }
    }

    // .....
    double avg=0;
    std::cout<<"Starting CPU computation"<<std::endl;
    for(int i=0; i <= iter; i++) {
        t=clock();
        CPUmatmul(N, x, y,ans);
        t = clock() - t;
        if(i)avg+=t; //we will ignore the first run
        // printf ("It took CPU-%d %f ms.\n",i,(((double)t)/CLOCKS_PER_SEC)*1000);
    }

    avg/=iter;
    avg/=CLOCKS_PER_SEC;
    avg*=1000;
    printf ("It took %lf ms on avg.\n",avg);
}

```

Matrix Multiplication on the CPU

- Sequential Matrix Multiplication
 - Here is a code for matrix multiplication using C++.
 - It is the standard $O(N^3)$ procedure.
- Here x, y, and ans are three N^2 size matrices
 - N^2 sized 1D array.
 - We are using 1D arrays as if it were 2D.

```

for(int i = 0; i < N; i++) {
    for(int j = 0; j < N; j++) {
        if(ans[i*N+j] != 20.0) return false;
    }
}
return true;
}

// ----- MAIN
int main(void)
{
    // size of matrix
    int N = 1<<9; // binary left-shift: 1 * 2^9 = 512
    printf("Size of matrix (N) is %d by %d.\n", N, N);
    int iter = 3;
    clock_t t;

    // Matrices
    double *x, *y, *ans;

    // TODO: Allocate Unified Memory - accessible from both CPU and GPU
    // ...
    // ...
    // ...

    // .....
    // initialize x,y and ans arrays on the host
    for (int i = 0; i < N; i++) {
        for(int j = 0; j < N; j++) {
            x[i*N+j] = 5;
            y[i*N+j] = (i==j?1:0);
            ans[i*N+j] = (double)0.000000000000;
        }
    }

    // .....
    double avg=0;
    std::cout<<"Starting unoptimized GPU computation"<<std::endl;
    // Run kernel on GPU
    for(int i = 0; i <= iter; i++) {
        t = clock();
        GPUmatmul<<<1,1>>>(N, x, y,ans);
        cudaDeviceSynchronize();
        t = clock() - t;
        if(i) avg += t; //we will ignore the first run
        // printf ("It took GPU-%d %f ms.\n",i,(((double)t)/CLOCKS_PER_SEC)*1000);
    }

    avg /= iter;
    avg /= CLOCKS_PER_SEC;
    avg *= 1000;
    printf("It took %lf ms on avg.\n", avg);
    if(check(N,ans)) std::cout<<"RUN OK."<<std::endl;
    else std::cout<<"RUN NOT OK."<<std::endl;

    // .....

    // TODO: Free memory
    // ...

```

Matrix Multiplication on the GPU

- This is the code to run on the GPU.
 - But it only uses one GPU thread.
 - And it is still sequential.
- You are tasked to use a simple stride pattern to make this parallel on the GPU.

Source Code

- Code from the lecture and project:

https://github.com/mertside/CS5375_GPU_Lecture



Readings

- How to CUDA? GPU Accelerated Computing with C and C++:
 - <https://developer.nvidia.com/how-to-cuda-c-cpp>
- Introduction to CUDA:
 - <https://developer.nvidia.com/blog/even-easier-introduction-cuda/>
- Unified Memory with CUDA:
 - <https://developer.nvidia.com/blog/unified-memory-cuda-beginners/>
- How to Optimize Data Transfers in CUDA C/C++:
 - <https://developer.nvidia.com/blog/how-optimize-data-transfers-cuda-cc/>
- An Efficient Matrix Transpose in CUDA using Shared Memory:
 - <https://developer.nvidia.com/blog/efficient-matrix-transpose-cuda-cc/>