



# **CS5375 Computer Systems Organization and Architecture**

## **Lecture 20**

Instructor: Yong Chen, Ph.D.  
Department of Computer Science  
Texas Tech University  
Yong.Chen@ttu.edu, 806-834-0284

## Announcements

- HW#4 posted and due on 11/17, Thursday
  - Note there are 8 questions in total (not 10 questions as in the prior homework)
- Conference trip from this afternoon till early next week
  - Class on 11/15, Tue., is skipped, leave the time for you to work on HW#4 or PP#2
- Will reduce one homework to have five homework in total, instead of six
  - Written assignments – 30% (~~six~~ five written assignments, weighing ~~5%~~ 6% each)
  - Last homework, HW#5, will be for covered Chapter 5 and Chapter 6 materials

## Outline

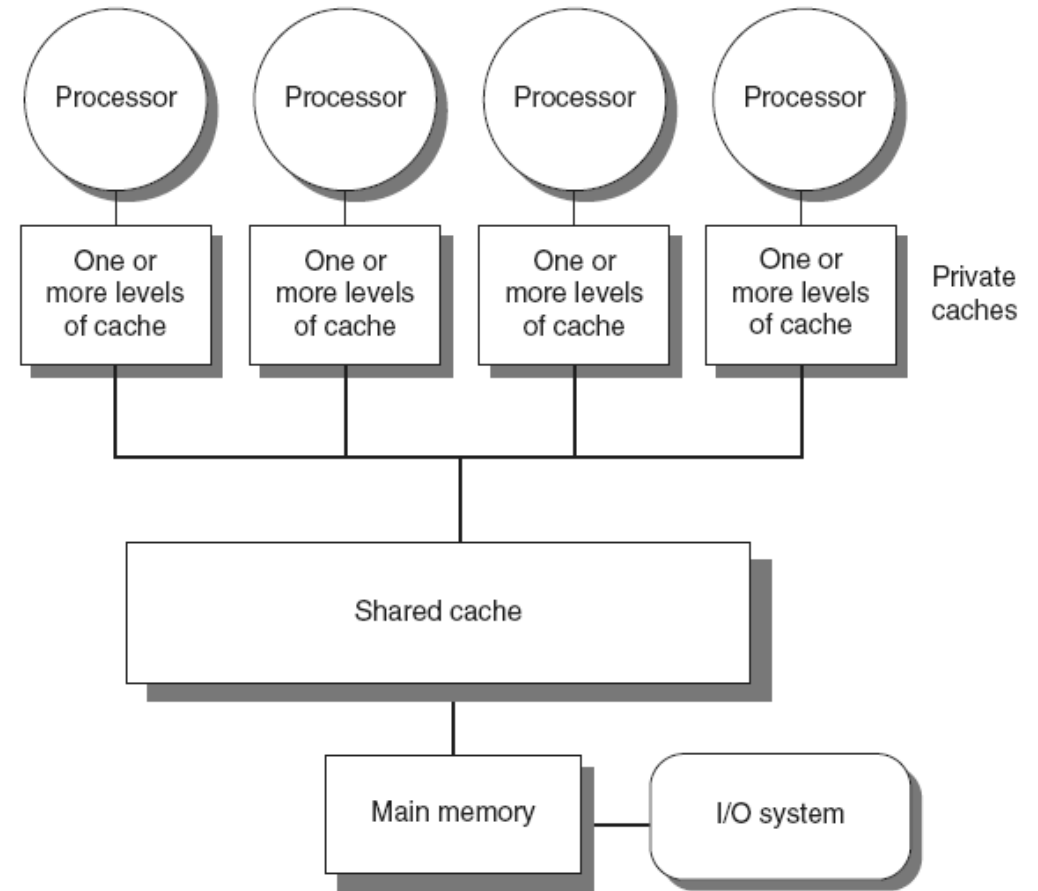
- Thread-Level Parallelism

# Introduction

- Thread-Level Parallelism (TLP)
  - Have multiple program counters
  - Uses MIMD model
    - Multiple Instruction: every processor may execute different instruction stream, may have separate clocks
    - Multiple Data: every processor may be working with a different data stream
  - Targeted for tightly-coupled **shared-memory multiprocessors**
    - V.S. **distributed-memory multicomputers** (“clusters”, “warehouse-scale computers”)
- Amount of computation assigned to each thread = grain size
  - Threads can be used for data-level parallelism, but the overheads may outweigh the benefit

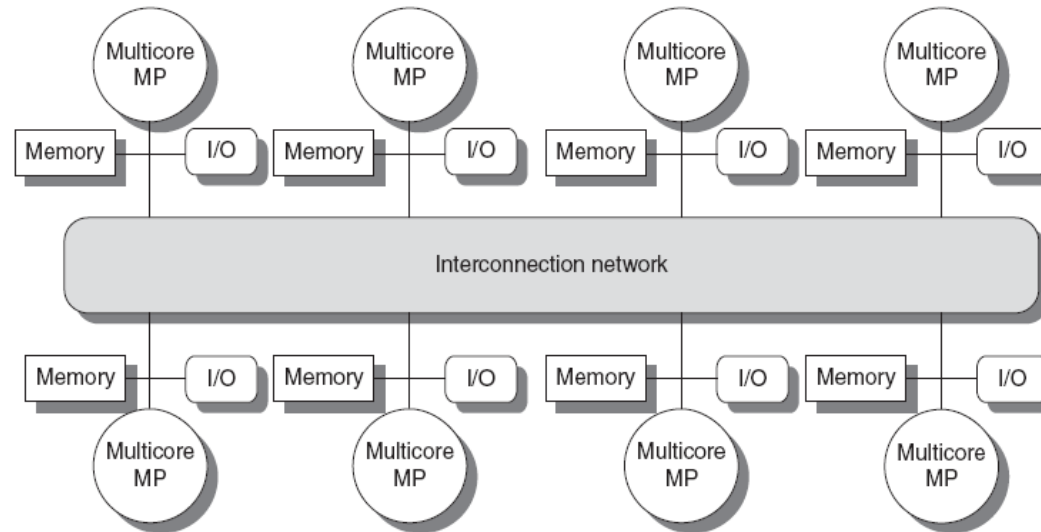
## Types

- Symmetric multiprocessors (SMP)
  - Small number of cores
  - Share single memory with uniform memory latency
- Also called as centralized shared-memory multiprocessors



## Types (cont.)

- Distributed shared memory (DSM)
  - Memory distributed among processors
  - Non-uniform memory access/latency (NUMA)
  - Processors connected via direct (switched) and non-direct (multi-hop) interconnection networks



- Note this is **different from “distributed-memory multicomputers”**, where memory is not shared, and some form of “message passing” to solve a problem

## Cache Coherence

- Processors may see different values through their caches (private caches):


Time	Event	Cache contents for processor A	Cache contents for processor B	Memory contents for location X
0				1
1	Processor A reads X	1		1
2	Processor B reads X	1	1	1
3	Processor A stores 0 into X	0	1	0

Assume a **write-through** cache

- Such a problem is referred to as the **cache coherence** problem

# Cache Coherence


- P1 loads X from main memory into its cache





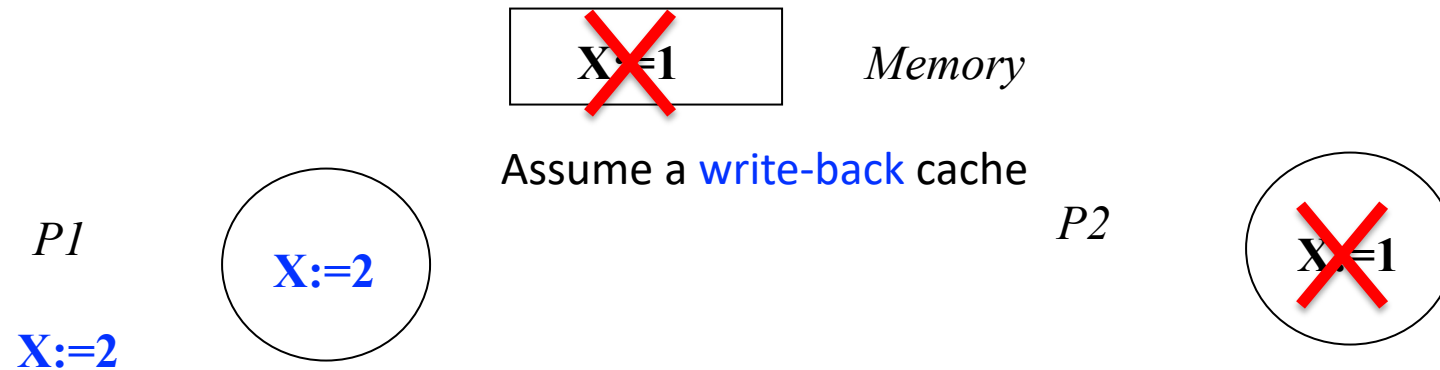
# Cache Coherence

- P1 loads X from main memory into its cache
- P2 loads X from main memory into its cache



# Cache Coherence

- P1 stores 2 into X
- We don't have consistent values for X across the memory hierarchy

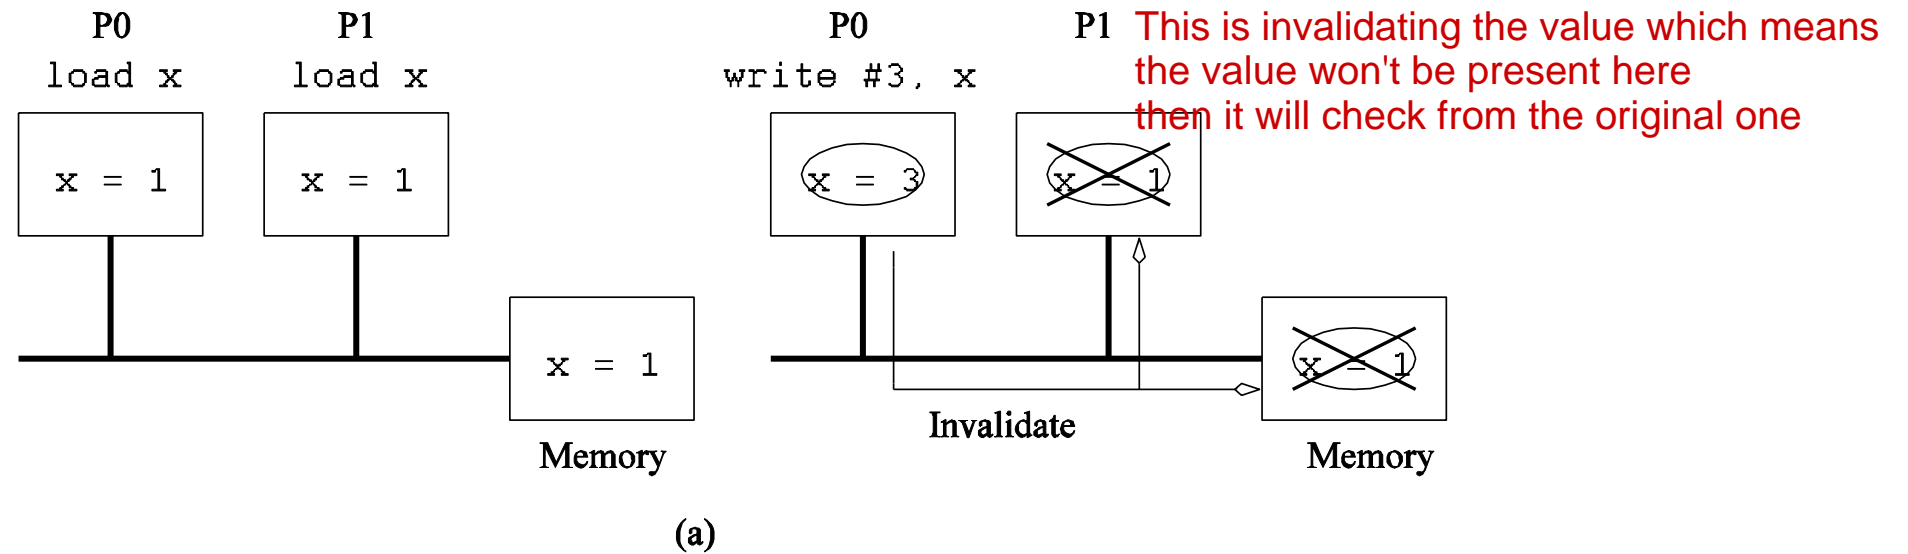


# Cache Coherence

- Cache coherence problem has two aspects
- **Coherence aspect:** defines *Single variable updation* **what** a written value returned by a read
  - All reads by any processor must return the most recently written value
  - Writes to the same location by any two processors are seen in the same order by all processors
- **Consistency aspect:** defines *This is for Multiple variables updation* **when** a written value returned by a read
  - If a processor writes location A followed by location B, any processor that sees the new value of B must also see the new value of A

## Enforcing Coherence

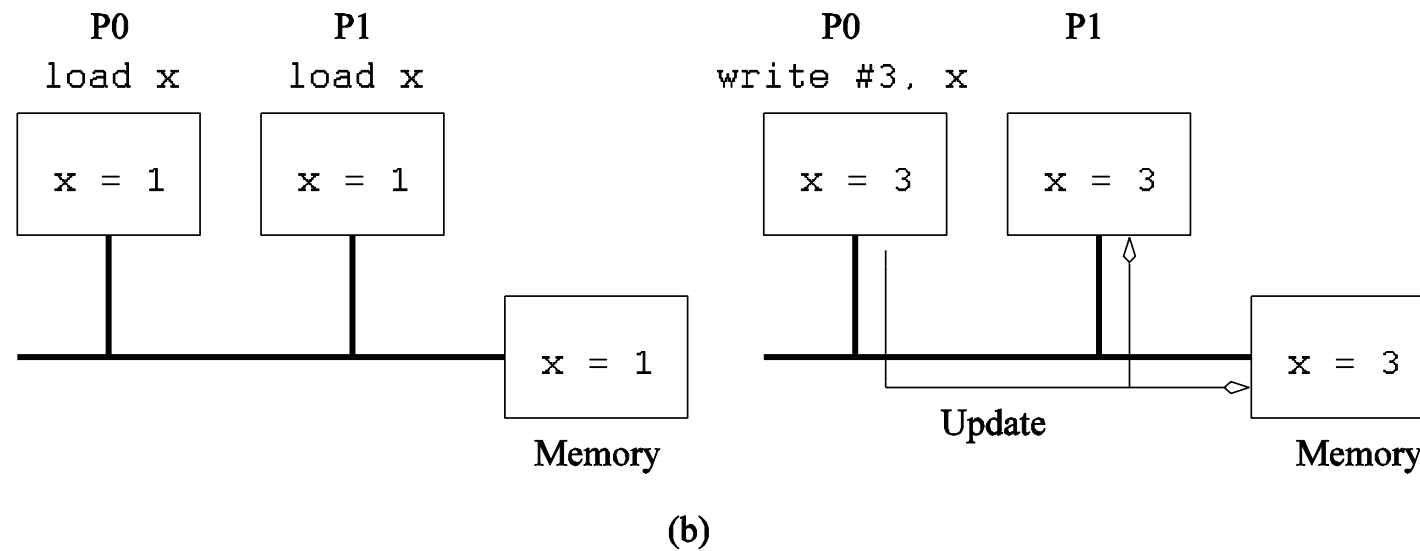
When the value of a variable is changed, all its copies must either be invalidated or updated.



Cache coherence in multiprocessor systems: (a) **Invalidate protocol**

## Enforcing Coherence (cont.)

When the value of a variable is changed, all its copies must either be invalidated or updated.



Cache coherence in multiprocessor systems: (b) **Update protocol** for shared variables.

## Update v.s. Invalidate Protocols

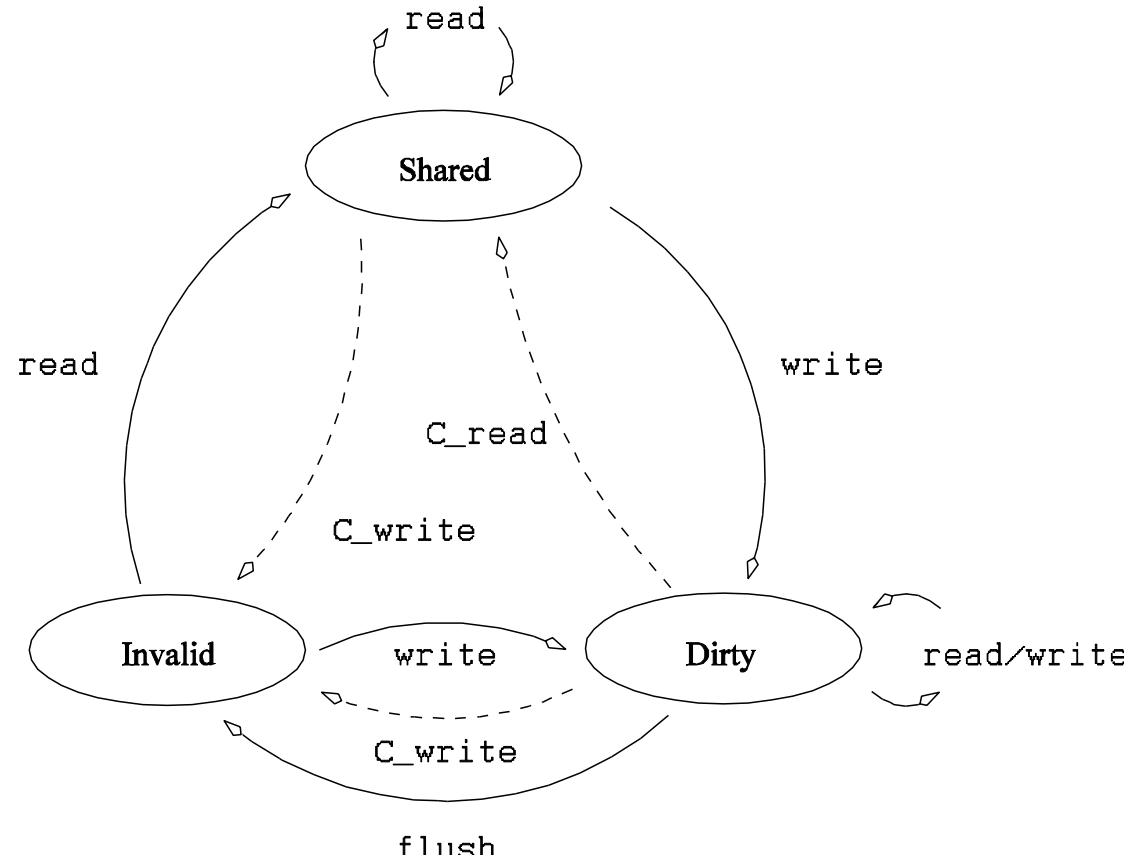
- If a processor just reads a value once and does not need it again, an update protocol may generate significant overhead
- If two processors make interleaved test and updates to a variable, an update protocol is better
  - Ping-pong
- Both protocols suffer from **false sharing** overheads (two words that are not shared, however, they lie on the same cache line)

it is a problem caused by cahce
- Most current machines use invalidate protocols

## Using Invalidate Protocols

- Each copy of a data item is associated with a state
- One example of such a set of states is, shared, invalid, or dirty/exclusive
- In shared state, there are multiple valid copies of the data item (and therefore, an invalidate would have to be generated on a write (or store) )
- In dirty state, only one copy exists and therefore, no invalidates need to be generated
- In invalid state, the data copy is invalid, therefore, a read (or load) generates a data request (and associated state changes)

## Using Invalidate Protocols (cont.)



State diagram of a simple three-state coherence protocol.



## Using Invalidate Protocols (cont.)


Time ↓	Instruction at Processor 0	Instruction at Processor 1	Variables and their states at Processor 0	Variables and their states at Processor 1	Variables and their states in Global mem.	
					x = 5, D y = 12, D	D for dirty
	read x		x = 5, S		x = 5, S	
		read y		y = 12, S	y = 12, S	S for shared
	x = x + 1		x = 6, D		x = 5, I	I for invalid
		y = y + 1		y = 13, D	y = 12, I	
	read y		y = 13, S	y = 13, S	y = 13, S	
		read x	x = 6, S	x = 6, S	x = 6, S	
	x = x + y		x = 19, D	x = 6, I	x = 6, I	
		y = x + y	y = 13, I	y = 19, D	y = 13, I	
	x = x + 1		x = 20, D		x = 6, I	
		y = y + 1		y = 20, D	y = 13, I	

Example of a thread program execution with the simple three-state coherence protocol

# Snoopy/Snooping Cache Systems

How are invalidates sent to the right processors?

In snoopy caches, there is a broadcast media that listens to all invalidates and read requests and performs appropriate coherence operations locally



A simple snooping bus based cache coherence system.

## Performance of Snoopy Caches

- Once copies of data are tagged dirty, all subsequent operations can be performed locally on the cache without generating external traffic.
- If a data item is read by a number of processors, it transitions to the shared state in the cache and all subsequent read operations become local
- If processors read and update data at the same time, they generate coherence requests on the bus - which is ultimately bandwidth limited

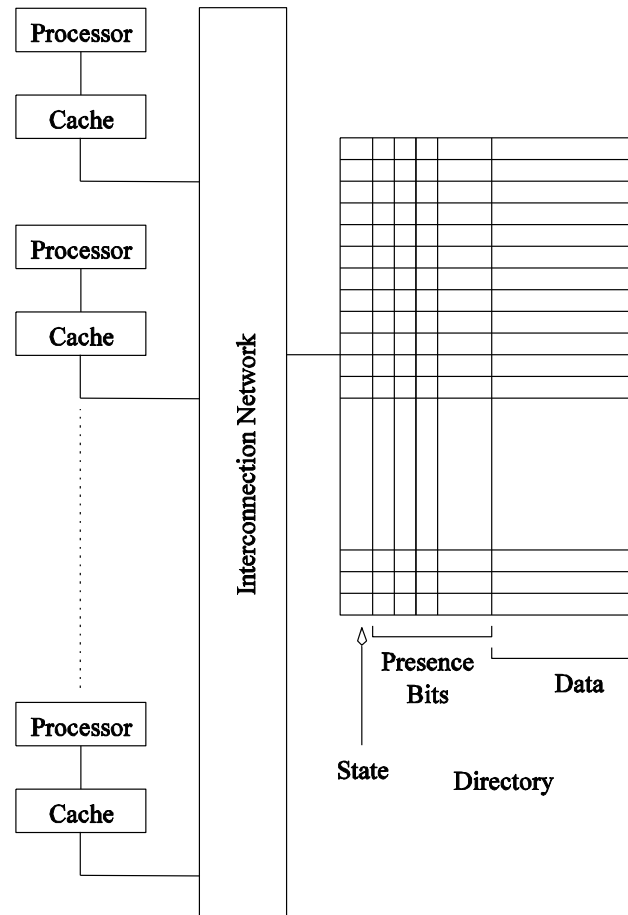
## Performance of Snoopy Caches

- Coherence influences cache miss rate
  - Coherence misses
  - True sharing misses
    - Write to shared block (transmission of invalidation)
    - Read an invalidated block
  - False sharing misses
    - Read an unmodified word in an invalidated block

## Directory Based Cache Coherence Systems

- In snoopy caches, each coherence operation is sent to all processors
  - This is an inherent limitation
- Why not send coherence requests to only those processors that need to be notified?
- This is done using a directory, which maintains a presence vector for each data item (cache line) along with its state.

# Directory Based Cache Coherence Systems



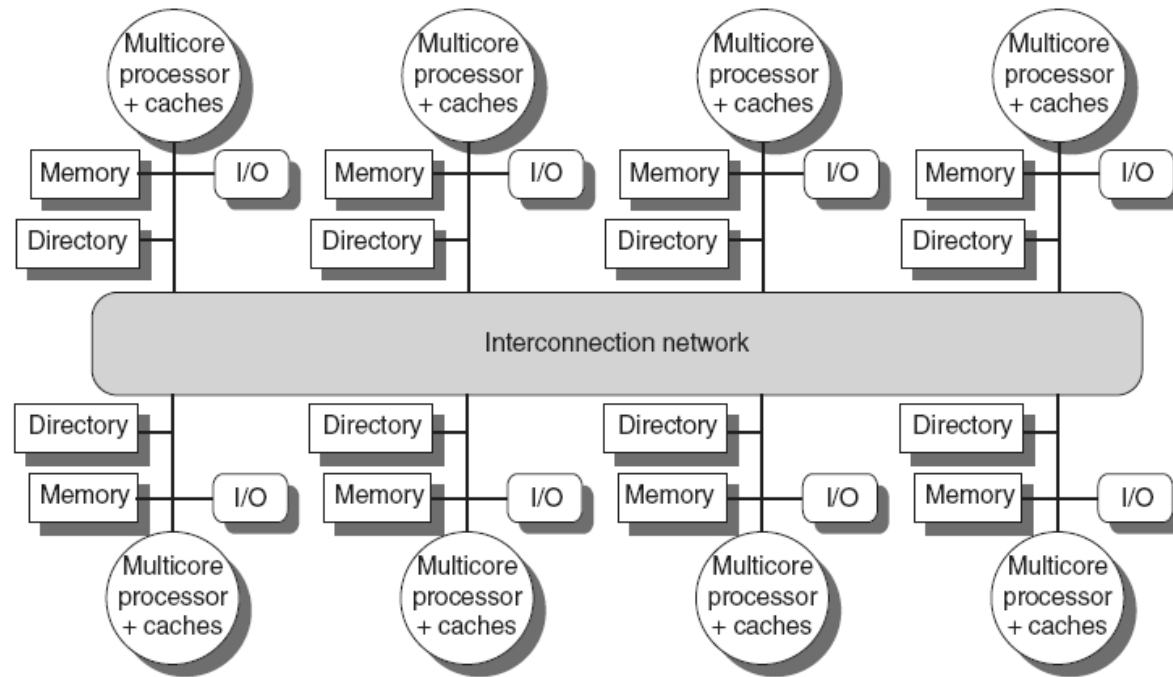
Architecture of typical directory based systems: a centralized directory

## Performance of Directory Based Cache Coherence

- The need for a snoopy hardware is replaced by the directory
- The additional bits to store the directory may add significant overhead
  - How much overhead?
- The interconnection network must be able to carry all the coherence requests
- The **directory is a point of contention**, therefore, distributed directory schemes can be used

# Directory Based Cache Coherence Systems

- Permit  $O(p)$  simultaneous coherence operations
- More scalable than snoopy or centralized directory systems
- Remains significant overhead of directory storage



Architecture of typical directory based systems: a distributed directory



## Readings

- Chapter 5, 5.1-5.4