




CS5375 Computer Systems Organization and Architecture

Lecture 13

Instructor: Yong Chen, Ph.D.
Department of Computer Science
Texas Tech University
Yong.Chen@ttu.edu, 806-834-0284

Review of Last Lecture

- Instruction-level parallelism (ILP) and Its Exploitation
- Control Hazards and Branch Prediction
 - Even with added hardware to compare registers, compute target and update PC in the ID stage, there is still a pipeline stall for every conditional branch (~10%) – need additional solution
- The solution is to predict the outcome of branch
 - Static branch prediction
 - Always predicts taken or not taken
 - Dynamic branch prediction
 - Hardware measures actual branch behavior (recent history) and predicts
 - Works because underlying algorithms and data being operated have regularities
 - 1-bit predictor, 2-bit predictor
 - Correlating predictors: record m most recently executed branches as taken / not taken and use that pattern to select proper n-bit branch history table, (m,n) predictor
 - Tournament predictors: dynamically choose between global and local predictors



Outline

- Exploiting ILP Using Multiple Issue and Static Scheduling
- Dynamic Scheduling

Multiple issue and Static Scheduling

- To further increase ILP
 - Multiple issue (N-way multi-issue CPU)
 - Replicate pipeline stages \Rightarrow multiple pipelines (i.e., multiple washers, dryers, “folders”, “storsers” in the laundry example)
 - Start multiple instructions per clock cycle
 - Launching multiple instructions per stage allows the instruction execution rate to exceed the clock rate, or in other words, $CPI < 1$
 - Thus it's also often to use Instructions Per Cycle (IPC), the inverse of CPI
 - E.g., 4-way multiple-issue 4GHz CPU
 - Peak $CPI = 0.25$, peak $IPC = 4$, 16 BIPS (billion instructions per second)
 - But hazards (dependencies) reduce this in practice

Multiple Issue

- **Static multiple issue** (at compile time, “static”)
 - Compiler groups instructions to be issued together
 - Packages them into “issue slots”
 - Compiler detects and avoids hazards
- **Dynamic multiple issue** (during execution, “dynamic”)
 - CPU examines instruction stream and chooses instructions to issue each cycle
 - Compiler can help by reordering instructions
 - CPU resolves hazards using advanced techniques at runtime

Static Multiple Issue

- Compiler groups instructions into “issue packets”
 - Issue packet: a set of instructions that can be issued together in one clock cycle
 - E.g., an arithmetic instruction and a load/store instruction, if there's no dependence/no hazards
- An issue packet can be considered as a very long instruction
 - Specifies multiple concurrent operations
 - ⇒ Very Long Instruction Word (VLIW)
 - VLIW also refers to a style of ISA that launches many operations that are defined to be independent in a single wide instruction (typically with many separate opcode fields)

RISC-V with Static Dual Issue (Two-way/Two-issue)

- Two-issue packets
 - One ALU/branch instruction
 - One load/store instruction
 - Fetching and decoding 64 bits of instructions (two instruction words)
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

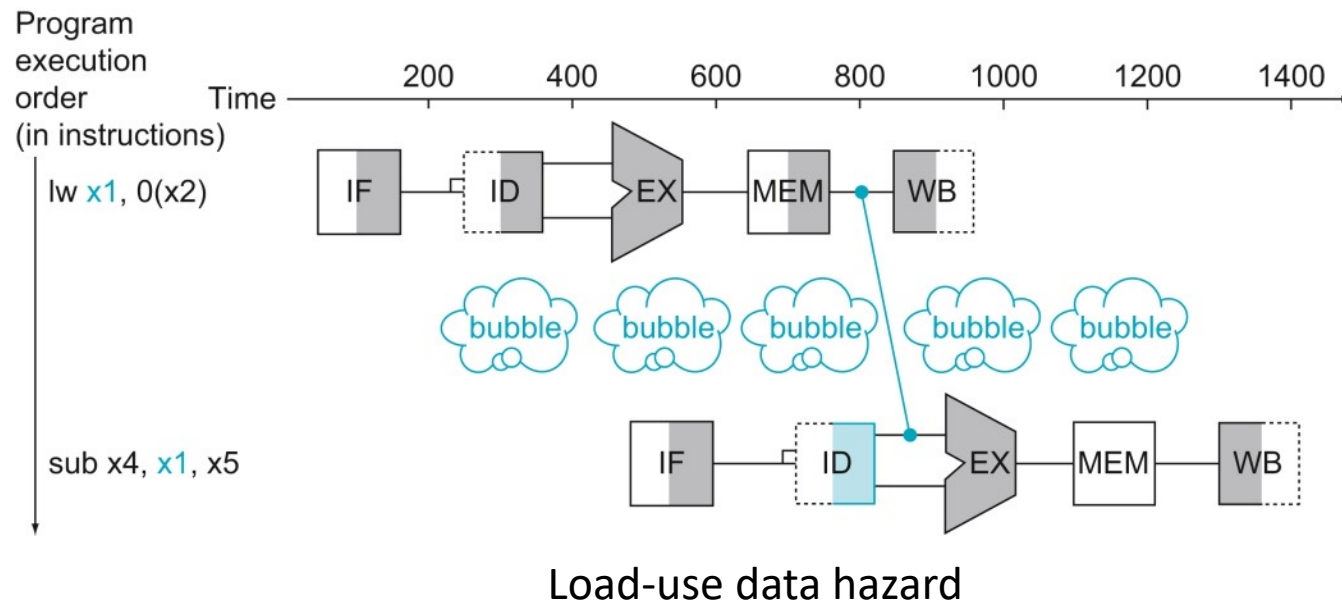
Two-issue pipeline example (v.s. single-issue pipeline in the earlier lecture)

Hazards in the Dual-Issue RISC-V

- Ideally two-issue processor can improve performance by up to a factor of two
- Would require more instructions executing in parallel
 - This increases the relative performance loss from data and control hazards
- E.g., data hazard
 - Forwarding avoided stalls with single-issue
 - Now can't use ALU result in load/store in same packet
 - add x10, x0, x1
 - lw x2, 0(x10)
 - Split into two packets, effectively a stall

Hazards in the Dual-Issue RISC-V (cont.)

- Another example: load-use hazard
 - A load has a **use latency** of one clock cycle, which means an instruction that can use the result of the load has to be one cycle later
 - Thus, in a two-issue processor, **next two instructions cannot use the load result without stalling**
 - In general, the speedup is less than 2 for a two-issue processor
 - More aggressive scheduling required



Compiler Scheduling Example

- Schedule this for dual-issue RISC-V

```
for (i=99; i>=0; i--)  
    a[i] = a[i] + s;  
  
Loop: lw    x31, 0(x20)      // x31=array element  
      add   x31, x31, x21    // add scalar in x21  
      sw    x31, 0(x20)      // store result  
      addi  x20, x20, -4     // decrement pointer  
      blt   x22, x20, Loop   // branch if x22 < x20
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw x31, 0(x20)	1
	addi x20, x20, -4	nop	2
	add x31, x31, x21	nop	3
	blt x22, x20, Loop	sw x31, 4(x20)	4

CPI = 4 cycles/5 instructions = 0.8 (v.s. an ideal case of 0.5)

Or, IPC = 5/4 = 1.25 (v.s. a peak IPC = 2)

Loop Unrolling and Scheduling

- Replicate loop body to expose more parallelism
- Loop unrolling example

```
for (i=99; i>=0; i=i-4){  
    a[i] = a[i] + s;  
    a[i-1] = a[i-1] + s;  
    a[i-2] = a[i-2] + s;  
    a[i-3] = a[i-3] + s;  
}
```

Since x20 is decremented by 16 (unrolling 4 iterations), these addresses are the original value of x20 minus 4, minus 8, and minus 12

	ALU/branch	Load/store	cycle
Loop:	addi x20, x20, -16	lw x28, 0(x20)	1
	nop	lw x29, 12(x20)	2
	add x28, x28, x21	lw x30, 8(x20)	3
	add x29, x29, x21	lw x31, 4(x20)	4
	add x30, x30, x21	sw x28, 16(x20)	5
	add x31, x31, x21	sw x29, 12(x20)	6
	nop	sw x30, 8(x20)	7
	blt x22, x20, Loop	sw x31, 4(x20)	8

- IPC?

$$\text{IPC} = 14/8 = 1.75$$

Closer to 2, but at cost of registers and code size

Loop Unrolling and Scheduling (cont.)

- These sequences of instructions in this loop are actually completely independent (adding a scalar value to each element of an array)

```
Loop: lw    x31, 0(x20)      // x31=array element
      add   x31, x31, x21    // add scalar in x21
      sw    x31, 0(x20)     // store result
      addi  x20, x20, -4     // decrement pointer
      blt   x22, x20, Loop  // branch if x22 < x20
```

- Except using x31 register in each loop iteration, which causes an enforced ordering by the reuse of a name (x31 register)
 - Store followed by a load of the same register, i.e. WAR or antidependence (a name dependence)
- Loop unrolling uses register renaming, i.e. 4 temporary registers rather than one, to resolve antidependence
 - Significant increase in code size, but delivers better performance (IPC = 1.75 or CPI = 0.57)

Outline

- Exploiting ILP Using Multiple Issue and Static Scheduling
- Dynamic Scheduling

Dynamic Scheduling

- Rearrange order of instructions to reduce stalls while maintaining data flow
- Advantages:
 - Compiler doesn't need to have knowledge of microarchitecture to rearrange instructions
 - Handles cases where dependencies are unknown at compile time
- Disadvantages:
 - Substantial increase in hardware complexity
 - Complexity leads to the power wall

Dynamic Scheduling

- Dynamic scheduling implies:
 - Out-of-order (OOO) execution and completion

- Example 1:

`fdiv.d f0, f2, f4`

`fadd.d f10, f0, f8`

`fsub.d f12, f8, f14`

`fsub.d` is not dependent, can issue before `fadd.d`

Dynamic Scheduling

- Example 2:

I1: `fdiv.d f0,f2,f4`

I2: `fmul.d f6,f0,f8`

I3: `fadd.d f0,f10,f14`

I1 and I2: RAW dependence (true dependence)

I2 and I3: WAR dependence (antidependence)

I1 and I3: WAW dependence (output dependence)

`fadd.d` is not dependent, but the antidependence and output dependence make it impossible to issue earlier without register renaming (e.g., use `f12` instead of `f0`, as `fadd.d f12,f10,f14`)

Register Renaming

- Example 3:

<code>fdiv.d</code>	<code>f0, f2, f4</code>	
<code>fadd.d</code>	<code>f6, f0, f8</code>	
<code>fsd</code>	<code>f6, 0(x1)</code>	
<code>fsub.d</code>	<code>f8, f10, f14</code>	
<code>fmul.d</code>	<code>f6, f10, f8</code>	

Output dependence

Antidependence

Register Renaming

- Example 3:

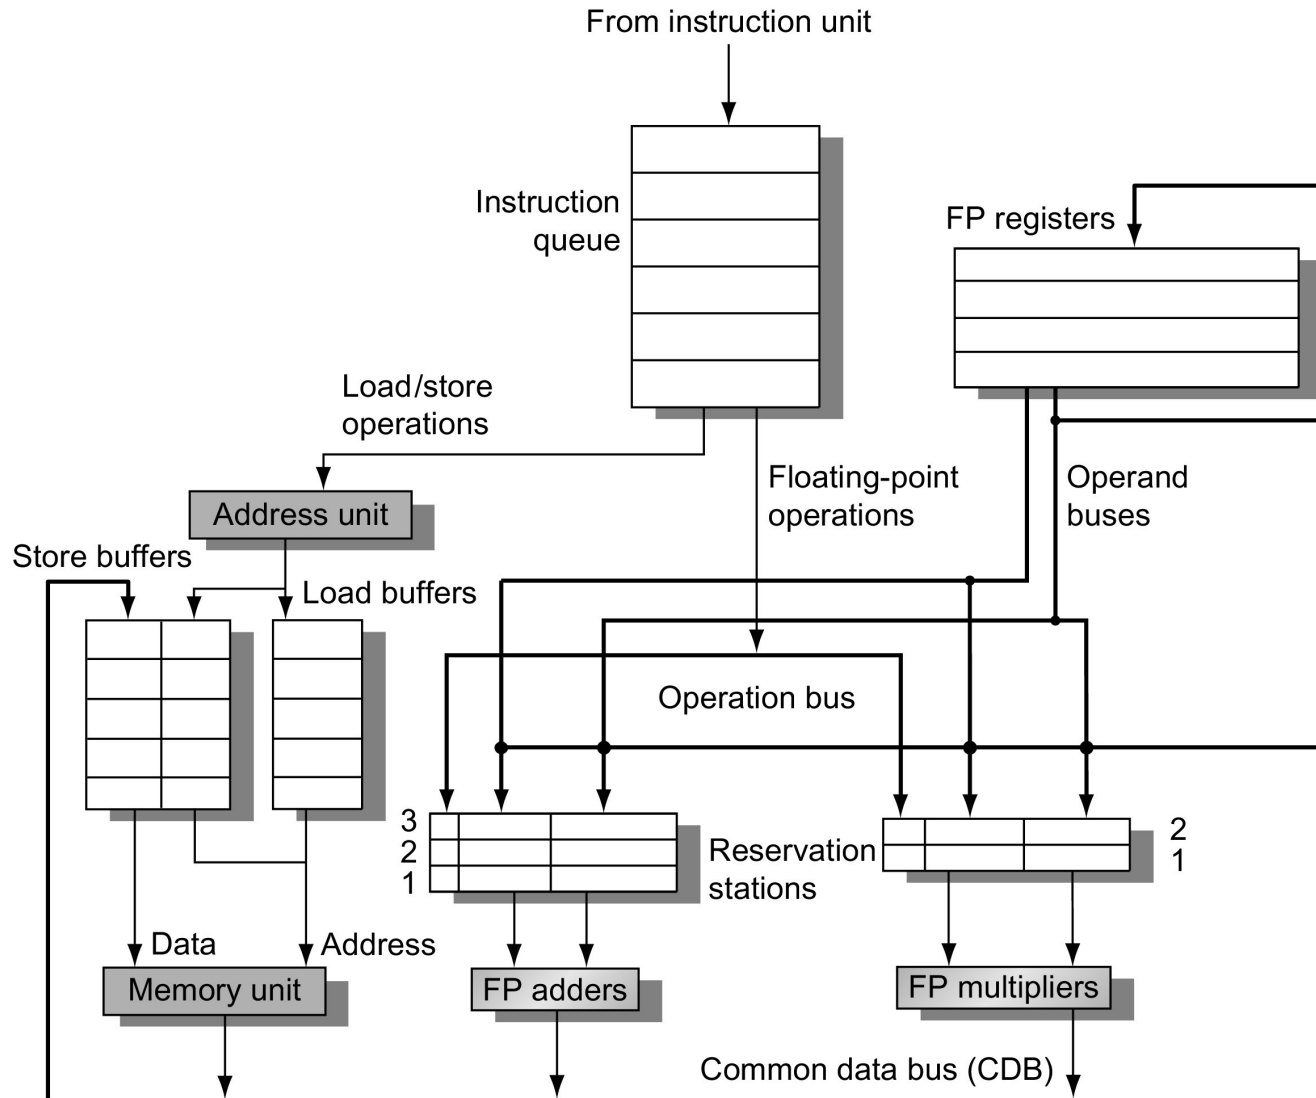
```
fdiv.d    f0, f2, f4
fadd.d    S, f0, f8
fsd       S, 0(x1)
fsub.d    T, f10, f14
fmul.d    f6, f10, T
```

- Now only RAW hazards remain

Dynamic Scheduling Using Tomasulo's Approach

- Tomasulo's Approach
 - A method of implementing dynamic scheduling, invented by Robert Tomasulo
 - Tracks when operands are available, satisfying the data dependence
 - Introduces register renaming in hardware to remove name dependence
 - Minimizes WAW and WAR hazards
- Register renaming is provided by reservation stations (RS)
- RS contains:
 - OP code: the operation of the instruction
 - Qj, Qk: RS producing the source operands S1 and S2
 - Vj, Vk: value of the source operands S1 and S2
 - A: the effective address for load/store instructions
 - Busy: indicates whether this RS is be used or available

Tomasulo's Algorithm



Tomasulo's Algorithm

- Three Steps:
 - Issue
 - Get next instruction from FIFO queue
 - If available RS, issue the instruction to the RS with operand values if available
 - If operand values not available, stall the instruction
 - Execute/dispatch
 - When operand becomes available, store it in any reservation stations waiting for it
 - When all operands are ready, issue the instruction
 - Loads and store maintained in program order through effective address
 - Write result
 - Write result on CDB into reservation stations and store buffers

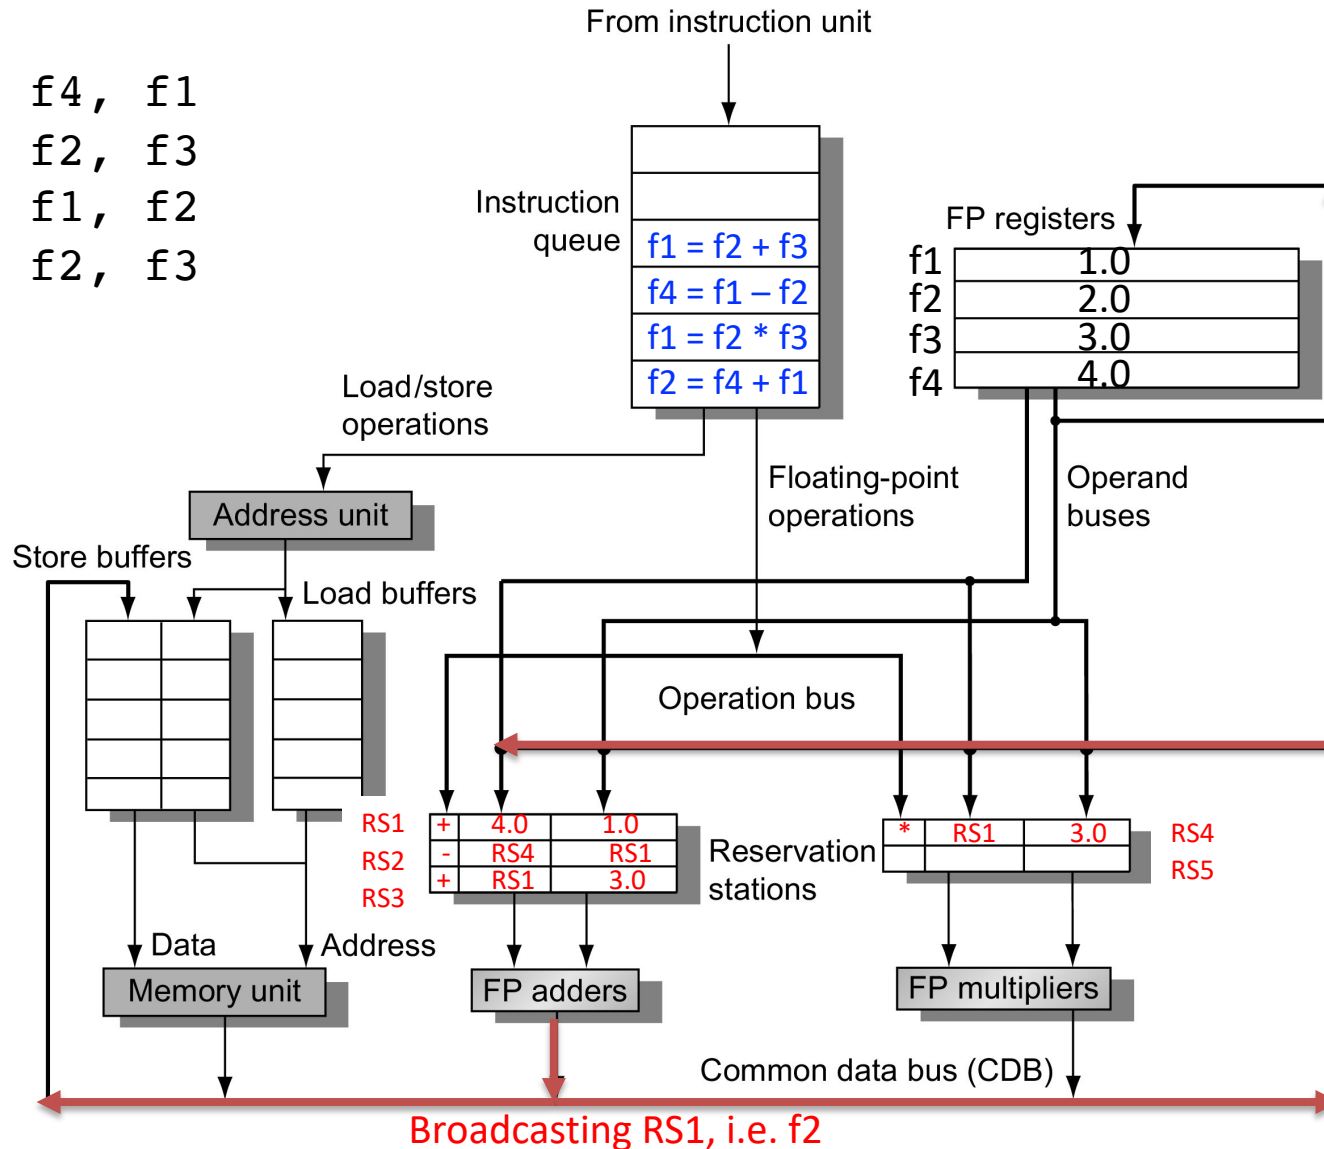
Tomasulo's Algorithm: An Example

I1: fadd.d f2, f4, f1
 I2: fmul.d f1, f2, f3
 I3: fsub.d f4, f1, f2
 I4: fadd.d f1, f2, f3

Register alias table

f1	RS3 RS4
f2	RS1
f3	
f4	RS2

If there's no RS3, we cannot issue I4.



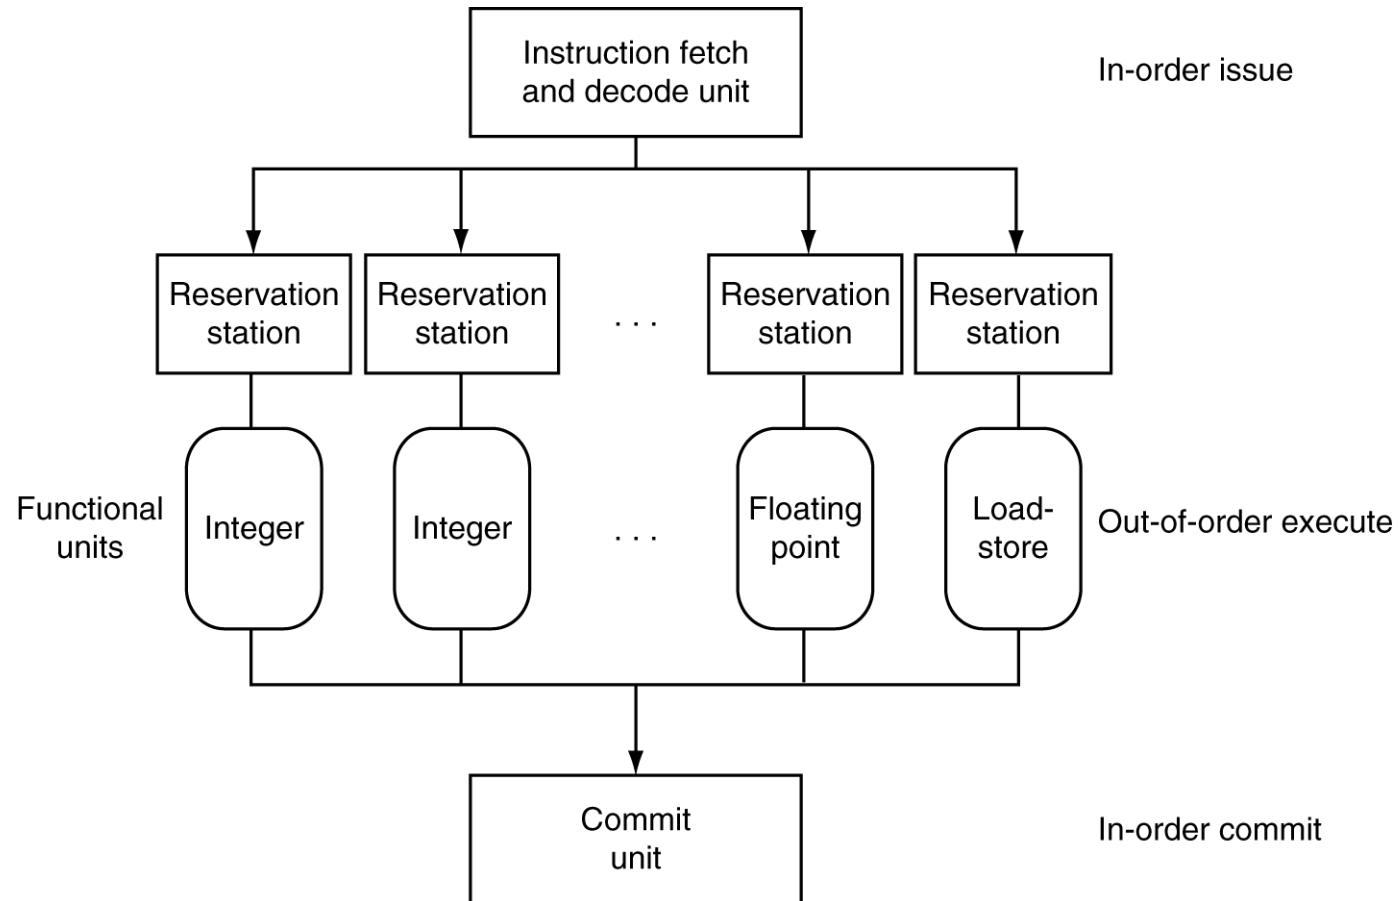
Such a process continues as issuing “in order” but executing “out of order”.

CPUs nowadays use similar mechanism to issue all instructions with 100s instructions in IQ.

Dynamic Multiple Issue

- Also known as “superscalar” processors, or “dynamic pipeline scheduling”
- CPU decides whether to issue 0, 1, 2, or more instructions each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling

Dynamic Multiple Issue CPU



Dynamic Scheduling, Multiple Issue, and Speculation

- Modern microarchitectures:
 - Dynamic scheduling + multiple issue + speculation
- How much to speculate
 - Mis-speculation degrades performance and power relative to no speculation
 - May cause additional misses (cache, TLB)
 - Prevent speculative code from causing higher costing misses (e.g. L2)
- Speculating through multiple branches
 - Complicates speculation recovery
- Speculation and energy efficiency
 - Note: speculation is only energy efficient when it significantly improves performance

Readings

- Chapter 3, 3.4 – 3.8